

Yuzhou Nie

Personal Website  yuzhounie@ucsb.edu  (+1)7656945966  rucnyz  Zhihu



Work on Post-training of Avocado. Side interest in inference infrastructure optimization.

Working Experience

Meta Superintelligence Lab
Senior AI Research Scientist

July.2026–Present

Education

University of California, Santa Barbara, CA, United States

Jan.2024–Dec.2026

Computer Science PhD, PhD Committee: Wenbo Guo, Dawn Song, Bo Li, Christopher Kruegel, Eric Xin Wang

Renmin University of China(RUC), Beijing, China

Sept.2019–June.2023

Bachelor of Science, School of Statistics **GPA:** GPA 3.85/4.00(Ranking: 3/52)

Awards

Anthropic Fellowship Program	August 2026
FAR AI PhD Fellowship	Supported from 2024 to 2025
UCSB Research Excellent Award	2024
Travel Grants for Different Conferences	2024–2026
National Scholarship	2020

Open Source Contributions

- **slime**: Added FSDP training support for Qwen3Next architecture (PR #1116).
- **flash-linear-attention**: Primary implementer of the Titans architecture (#159, #214); Other fixes (#687).
- **nitrobox** (Author): Lightweight Docker alternative with instant reset and process-level checkpointing.
- **prelude** (Author): Zero-dependency LLM inference engine in Rust; parity with vLLM/SGLang on Qwen.

Internship Experience

Virtue AI, San Francisco, CA
Founding Engineer

April.2025–

- Leading large-scale serving and post-training of cybersecurity-specialized LLMs.

Microsoft Research, Redmond, WA

June.2024–Sept.2024

Worked on data poisoning and backdoor attacks against large language models.

Amazon Web Services AI Lab, Shanghai, China

Aug.2022–March.2023

Internship, AWS Deep Learning group, Shanghai

- Graphical Structure for Electronic Health Records

Preprint

* indicates equal contribution

- **Yuzhou Nie**, Weifang Zhang, Zhen Xu, Songyang Peng, Yuheng Tang, Minzhou Pan, Bo Li, Dawn Song, Ce Zhang, Wenbo Guo, *HiMA: Efficient Hybrid Model Serving for Agentic Systems*, **ICML 2026 AIWild Workshop**.
- [ 300 ] **Yuzhou Nie**, Hongwei Li, Chengquan Guo, Ruizhe Jiang, Zhun Wang, Bo Li, Dawn Song, Wenbo Guo, *VulnLLM-R: Specialized Reasoning LLM with Agent Scaffold for Vulnerability Detection*. [**#2 most-downloaded model on HuggingFace**, Feb.2026]
- Kaijie Zhu*, **Yuzhou Nie***, Yijiang Li, Yiming Huang, Jialian Wu, Jiang Liu, Ximeng Sun, Zhenfei Yin, Lun Wang, Zicheng Liu, Emad Barsoum, William Yang Wang, Wenbo Guo, *TermiGen: High-Fidelity Environment and Robust Trajectory Synthesis for Terminal Agents*. [Submitted to ICML 2026; **HuggingFace Daily Paper #2**]

Publications

* indicates equal contribution

- **Yuzhou Nie***, Weifang Zhang*, Bowen Pang, Guangrui Ma, Shining Wu, *THETA: Threshold-Based Exclusive Batching for Memory-Bandwidth-Constrained LLM Inference*, **ICML 2026**.
- Chengquan Guo*, **Yuzhou Nie***, Chulin Xie, Zinan Lin, Wenbo Guo, Bo Li, *BlueCodeAgent: A Blue Teaming Agent Powered by Automated Red Teaming for CodeGen AI*, **ICML 2026**.

- Hongwei Li, Zhun Wang, Qinrun Dai, **Yuzhou Nie**, Jinjun Peng, Ruitong Liu, Jingyang Zhang, Kaijie Zhu, Jingxuan He, Lun Wang, Yangruibo Ding, Yueqi Chen, Wenbo Guo, Dawn Song, *OpenSage: Self-programming Agent Generation Engine*, **ICML 2026**.
- **Yuzhou Nie***, Zhun Wang*, Yu Yang*, Ruizhe Jiang, Yuheng Tang, Xander Davies, Yarin Gal, Bo Li, Wenbo Guo, Dawn Song, *SeCodePLT: A Unified Platform for Evaluating the Security of Code GenAI*, **NeurIPS 2025**.
- Mengkang Hu, Yuhang Zhou, Wendong Fan, **Yuzhou Nie**, Bowei Xia, Tao Sun, Ziyu Ye, Zhaoxuan Jin, Yingru Li, Qiguang Chen, Zeyu Zhang, Yifeng Wang, Qianshuo Ye, Bernard Ghanem, Ping Luo, Guohao Li, *Owl: Optimized workforce learning for general multi-agent assistance in real-world task automation*, **NeurIPS 2025**.
- **Yuzhou Nie**, Zhun Wang, Ye Yu, Xian Wu, Xuandong Zhao, Nathaniel D. Bastian, Wenbo Guo, Dawn Song, *LeakAgent: RL-based Red-teaming Agent for LLM Privacy Leakage*, **COLM 2025**. [1st place, **Safety Track, Berkeley LLM Agents Hackathon**]
- Zhun Wang, Vincent Siu, Zhe Ye, Tianneng Shi, **Yuzhou Nie**, Xuandong Zhao, Chenguang Wang, Wenbo Guo, Dawn Song, *AgentVigil: Generic Black-Box Red-teaming for Indirect Prompt Injection against LLM Agents*, **EMNLP 2025**.
- Xuan Chen, **Yuzhou Nie**, Lu Yan, Yunshu Mao, Wenbo Guo, Xiangyu Zhang, *RL-JACK: Reinforcement Learning-powered Black-box Jailbreaking Attack against LLMs*, **TMLR 2026**.
- **Yuzhou Nie***, Xuan Chen*, Wenbo Guo, Xiangyu Zhang, *When LLM Meets DRL: Advancing Jailbreaking Efficiency via DRL-guided Search*, **NeurIPS 2024**.
- **Yuzhou Nie***, Fengjiao Gong* and Hongteng Xu, *Gromov-Wasserstein Multi-modal Alignment and Clustering*, **CIKM 2022**.

Skills

Proficient: Python (PyTorch, Scikit-learn, etc.), C/C++, Java, SQL, Git/Terminal, LaTeX, R

Familiar: MATLAB, PySpark, Linux/Unix